

assignment 1

January 9, 2022

1 MSDS 422 Assignment 1

The dependent variable of interest is house prices in Ames, Iowa ('SalePrice')

1.1 Data preparation

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: #import data
train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')
```

```
[3]: train.head()
```

```
[3]:
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

```
[4]: train.tail()
```

```
[4]:      Id  MSSubClass MSZoning  LotFrontage  LotArea Street Alley LotShape  \
1455  1456          60      RL          62.0    7917   Pave   NaN     Reg
1456  1457          20      RL          85.0   13175   Pave   NaN     Reg
1457  1458          70      RL          66.0    9042   Pave   NaN     Reg
1458  1459          20      RL          68.0    9717   Pave   NaN     Reg
1459  1460          20      RL          75.0    9937   Pave   NaN     Reg

      LandContour Utilities  ... PoolArea PoolQC  Fence MiscFeature MiscVal  \
1455          Lvl   AllPub  ...      0    NaN   NaN          NaN      0
1456          Lvl   AllPub  ...      0    NaN MnPrv          NaN      0
1457          Lvl   AllPub  ...      0    NaN GdPrv          Shed    2500
1458          Lvl   AllPub  ...      0    NaN   NaN          NaN      0
1459          Lvl   AllPub  ...      0    NaN   NaN          NaN      0

      MoSold YrSold  SaleType  SaleCondition  SalePrice
1455      8    2007         WD         Normal    175000
1456      2    2010         WD         Normal    210000
1457      5    2010         WD         Normal    266500
1458      4    2010         WD         Normal    142125
1459      6    2008         WD         Normal    147500
```

[5 rows x 81 columns]

```
[5]: train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null  int64
1   MSSubClass            1460 non-null  int64
2   MSZoning              1460 non-null  object
3   LotFrontage          1201 non-null  float64
4   LotArea              1460 non-null  int64
5   Street               1460 non-null  object
6   Alley               91 non-null    object
7   LotShape             1460 non-null  object
8   LandContour          1460 non-null  object
9   Utilities            1460 non-null  object
10  LotConfig            1460 non-null  object
11  LandSlope            1460 non-null  object
12  Neighborhood          1460 non-null  object
```

13	Condition1	1460	non-null	object
14	Condition2	1460	non-null	object
15	BldgType	1460	non-null	object
16	HouseStyle	1460	non-null	object
17	OverallQual	1460	non-null	int64
18	OverallCond	1460	non-null	int64
19	YearBuilt	1460	non-null	int64
20	YearRemodAdd	1460	non-null	int64
21	RoofStyle	1460	non-null	object
22	RoofMatl	1460	non-null	object
23	Exterior1st	1460	non-null	object
24	Exterior2nd	1460	non-null	object
25	MasVnrType	1452	non-null	object
26	MasVnrArea	1452	non-null	float64
27	ExterQual	1460	non-null	object
28	ExterCond	1460	non-null	object
29	Foundation	1460	non-null	object
30	BsmtQual	1423	non-null	object
31	BsmtCond	1423	non-null	object
32	BsmtExposure	1422	non-null	object
33	BsmtFinType1	1423	non-null	object
34	BsmtFinSF1	1460	non-null	int64
35	BsmtFinType2	1422	non-null	object
36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64
53	KitchenQual	1460	non-null	object
54	TotRmsAbvGrd	1460	non-null	int64
55	Functional	1460	non-null	object
56	Fireplaces	1460	non-null	int64
57	FireplaceQu	770	non-null	object
58	GarageType	1379	non-null	object
59	GarageYrBlt	1379	non-null	float64
60	GarageFinish	1379	non-null	object

```

61 GarageCars      1460 non-null    int64
62 GarageArea      1460 non-null    int64
63 GarageQual      1379 non-null    object
64 GarageCond      1379 non-null    object
65 PavedDrive      1460 non-null    object
66 WoodDeckSF      1460 non-null    int64
67 OpenPorchSF     1460 non-null    int64
68 EnclosedPorch   1460 non-null    int64
69 3SsnPorch       1460 non-null    int64
70 ScreenPorch     1460 non-null    int64
71 PoolArea        1460 non-null    int64
72 PoolQC          7 non-null       object
73 Fence           281 non-null     object
74 MiscFeature     54 non-null      object
75 MiscVal         1460 non-null    int64
76 MoSold          1460 non-null    int64
77 YrSold          1460 non-null    int64
78 SaleType        1460 non-null    object
79 SaleCondition   1460 non-null    object
80 SalePrice       1460 non-null    int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

```
[6]: train.describe()
```

```

[6]:
count      Id      MSSubClass  LotFrontage      LotArea  OverallQual  \
count  1460.000000  1460.000000  1201.000000    1460.000000  1460.000000
mean    730.500000    56.897260    70.049958   10516.828082    6.099315
std     421.610009    42.300571    24.284752    9981.264932    1.382997
min       1.000000    20.000000    21.000000    1300.000000    1.000000
25%     365.750000    20.000000    59.000000    7553.500000    5.000000
50%     730.500000    50.000000    69.000000    9478.500000    6.000000
75%    1095.250000    70.000000    80.000000   11601.500000    7.000000
max    1460.000000   190.000000   313.000000  215245.000000   10.000000

      OverallCond   YearBuilt  YearRemodAdd  MasVnrArea  BsmtFinSF1  ...  \
count  1460.000000  1460.000000  1460.000000  1452.000000  1460.000000  ...
mean     5.575342   1971.267808   1984.865753   103.685262   443.639726  ...
std     1.112799    30.202904    20.645407   181.066207   456.098091  ...
min       1.000000  1872.000000   1950.000000     0.000000     0.000000  ...
25%       5.000000  1954.000000   1967.000000     0.000000     0.000000  ...
50%       5.000000  1973.000000   1994.000000     0.000000   383.500000  ...
75%       6.000000  2000.000000   2004.000000   166.000000   712.250000  ...
max       9.000000  2010.000000   2010.000000  1600.000000  5644.000000  ...

      WoodDeckSF  OpenPorchSF  EnclosedPorch   3SsnPorch  ScreenPorch  \
count  1460.000000  1460.000000    1460.000000  1460.000000  1460.000000

```

mean	94.244521	46.660274	21.954110	3.409589	15.060959
std	125.338794	66.256028	61.119149	29.317331	55.757415
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	25.000000	0.000000	0.000000	0.000000
75%	168.000000	68.000000	0.000000	0.000000	0.000000
max	857.000000	547.000000	552.000000	508.000000	480.000000

	PoolArea	MiscVal	MoSold	YrSold	SalePrice
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	2.758904	43.489041	6.321918	2007.815753	180921.195890
std	40.177307	496.123024	2.703626	1.328095	79442.502883
min	0.000000	0.000000	1.000000	2006.000000	34900.000000
25%	0.000000	0.000000	5.000000	2007.000000	129975.000000
50%	0.000000	0.000000	6.000000	2008.000000	163000.000000
75%	0.000000	0.000000	8.000000	2009.000000	214000.000000
max	738.000000	15500.000000	12.000000	2010.000000	755000.000000

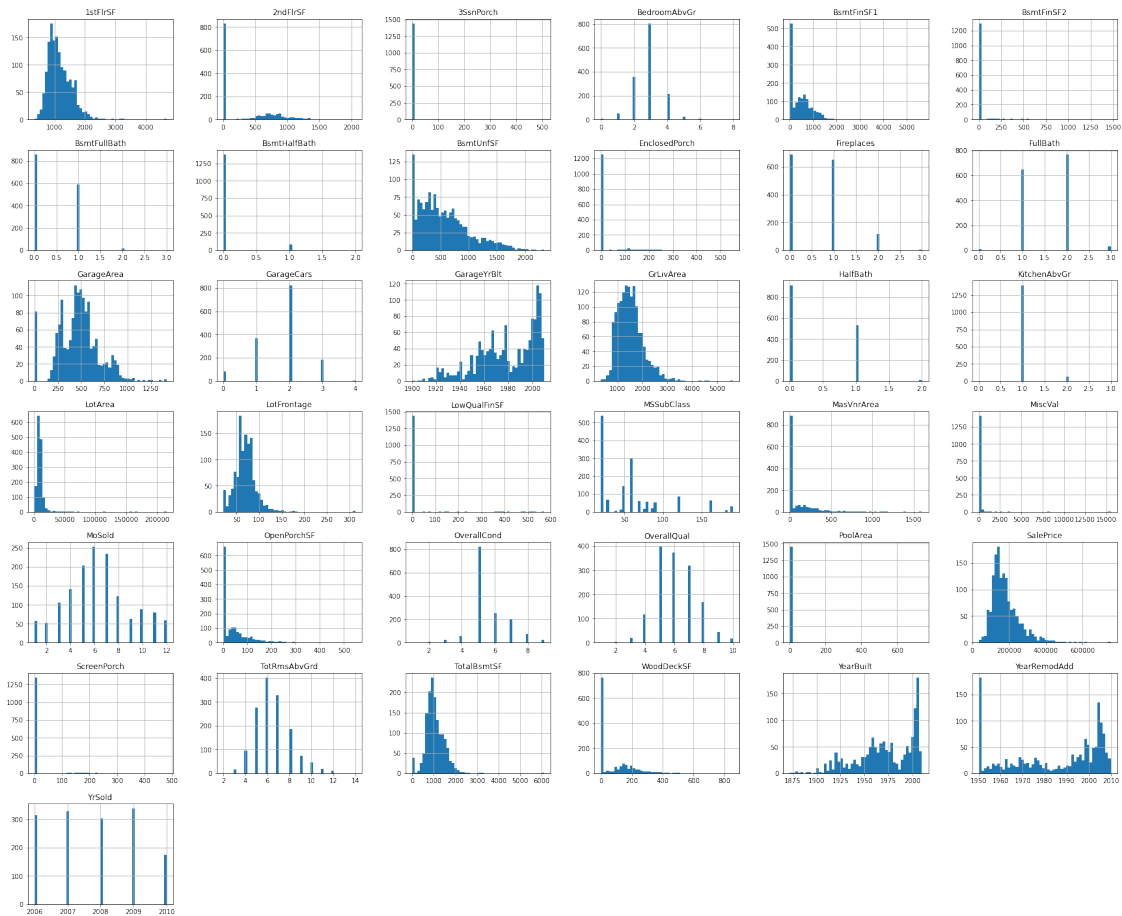
[8 rows x 38 columns]

```
[7]: train.shape
```

```
[7]: (1460, 81)
```

```
[8]: #no need for ID (arbitrary number)
train = train.drop(labels = 'Id',axis=1)
```

```
[9]: #distributions
train.hist(bins=50, figsize=(30,25))
plt.show()
```



1. Provide appropriate descriptive statistics and visualizations to help understand the marginal distribution of the dependent variable.


```
[10]: #code
train["SalePrice"].describe()
```

```
[10]: count      1460.000000
mean      180921.195890
std       79442.502883
min       34900.000000
25%      129975.000000
50%      163000.000000
75%      214000.000000
max       755000.000000
Name: SalePrice, dtype: float64
```

The average sale price from this data set is 180,921. Sale prices range from 129,975 to 755,000 with a median value of 163,000.

```
[11]: fig, ax = plt.subplots()
train["SalePrice"].hist(bins=50)
plt.xlabel('Sale Price ($)')
plt.ylabel('Counts')
plt.title('Histogram of Sale Prices')
```

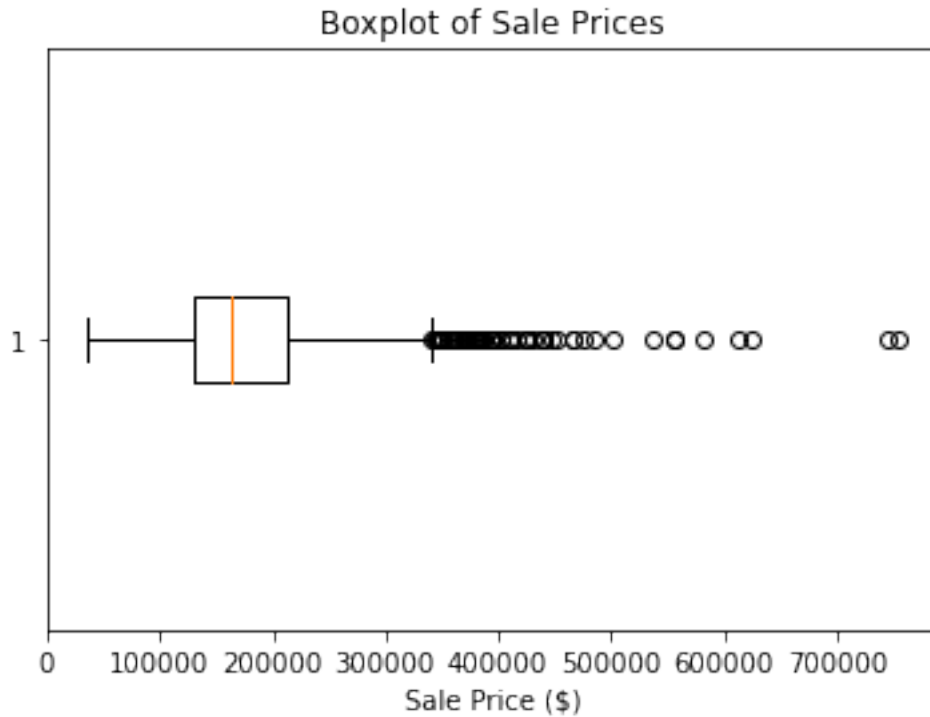
```
[11]: Text(0.5, 1.0, 'Histogram of Sale Prices')
```



```
[12]: fig, ax = plt.subplots()
plt.boxplot(train["SalePrice"], vert=False)
plt.xlabel('Sale Price ($)')
plt.title('Boxplot of Sale Prices')

#high end of boxplot
214000 + 1.5*(214000 - 129975)
```

```
[12]: 340037.5
```



The histogram and boxplot above show that SalePrice is skewed to the right. Outliers may range from any SalePrice greater than ~340,000

2. Investigate missing data and outliers.

```
[13]: #outliers
plt.figure(figsize = (140,20))
train.loc[:, train.columns!='SalePrice'].boxplot()
```

[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd8c4d18100>



```
[14]: #outliers
plt.figure(figsize = (140,20))
train.loc[:, ~train.columns.isin(['SalePrice', 'LotArea'])].boxplot()
```

[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd8c7552760>



SalePrice and LotArea have values with magnitudes greater than other independent variables. When looking at other variables, many of their boxplots indicate several outliers.

```
[15]: #missing data
train.isnull().sum()[train.isnull().sum()>0]
```

```
[15]: LotFrontage      259
      Alley          1369
      MasVnrType      8
      MasVnrArea      8
      BsmtQual        37
      BsmtCond        37
      BsmtExposure    38
      BsmtFinType1    37
      BsmtFinType2    38
      Electrical      1
      FireplaceQu     690
      GarageType      81
      GarageYrBlt     81
      GarageFinish    81
      GarageQual      81
      GarageCond      81
      PoolQC         1453
      Fence          1179
      MiscFeature     1406
      dtype: int64
```

The variables above have numerous outliers. Of which, Alley, PoolQC, Fence, and MiscFeature have over 1000 null values relative to the 1460 total records.

```
[16]: #duplicates
train.duplicated().sum()
```

```
[16]: 0
```

There are no duplicate rows in the dataset.

3. Investigate at least three potential predictors of the dependent variable and provide appropriate graphs / statistics to demonstrate the relationships.

```
[17]: #correlation matrix
train.corr().style.background_gradient(cmap='coolwarm').set_precision(2)
```

```
[17]: <pandas.io.formats.style.Styler at 0x7fd8cb2ccd30>
```

```
[18]: #highest correlations with SalePrice
train.corr()['SalePrice'].sort_values(ascending=False)
```

```
[18]: SalePrice      1.000000
OverallQual    0.790982
GrLivArea      0.708624
GarageCars     0.640409
GarageArea     0.623431
TotalBsmtSF    0.613581
1stFlrSF       0.605852
FullBath       0.560664
TotRmsAbvGrd   0.533723
YearBuilt      0.522897
YearRemodAdd    0.507101
GarageYrBlt     0.486362
MasVnrArea     0.477493
Fireplaces     0.466929
BsmtFinSF1     0.386420
LotFrontage    0.351799
WoodDeckSF     0.324413
2ndFlrSF       0.319334
OpenPorchSF    0.315856
HalfBath       0.284108
LotArea        0.263843
BsmtFullBath   0.227122
BsmtUnfSF      0.214479
BedroomAbvGr   0.168213
ScreenPorch    0.111447
PoolArea       0.092404
MoSold         0.046432
3SsnPorch      0.044584
BsmtFinSF2     -0.011378
BsmtHalfBath   -0.016844
MiscVal        -0.021190
LowQualFinSF   -0.025606
YrSold         -0.028923
OverallCond    -0.077856
MSSubClass     -0.084284
EnclosedPorch  -0.128578
KitchenAbvGr   -0.135907
Name: SalePrice, dtype: float64
```

The top 3 variables with the highest correlation with SalePrice are OverallQual (0.790982), GrLivArea (0.708624) and GarageCars (0.640409). All 3 have a positive correlation.


```
[19]: from sklearn.metrics import mean_squared_error
      from sklearn import datasets, linear_model
      from scipy import stats
```

```
[20]: lm_oq = linear_model.LinearRegression()
      lm_oq.fit(train[['SalePrice']], train['OverallQual'])
      x_oq = lm_oq.predict(train[['SalePrice']])
      rmse_oq = mean_squared_error(train['OverallQual'], x_oq, squared=False)

      plt.scatter(train['OverallQual'], train['SalePrice'], color='blue')
      plt.plot(x_oq, train['SalePrice'], color='maroon', linewidth=1)
      plt.ylabel('Sale Price ($)')
      plt.xlabel('Overall material and finish quality')
      plt.title('SalePrice vs OverallQual Regression')

      print("RMSE:", rmse_oq)
```

RMSE: 0.8458826841860243



```
[21]: lm_gla = linear_model.LinearRegression()
      lm_gla.fit(train[['SalePrice']], train['GrLivArea'])
```

```


x_gla = lm_gla.predict(train[['SalePrice']])
rmse_gla = mean_squared_error(train['GrLivArea'], x_gla, squared=False)

plt.scatter(train['GrLivArea'], train['SalePrice'], color='blue')
plt.plot(x_gla, train['SalePrice'], color='maroon', linewidth=1)
plt.ylabel('Sale Price ($)')
plt.xlabel('Above-ground living area (sqft)')
plt.title('SalePrice vs GrLivArea Regression')

print("RMSE:", rmse_gla)

```

RMSE: 370.6445090657199



```


[22]: lm_gc = linear_model.LinearRegression()
lm_gc.fit(train[['SalePrice']], train['GarageCars'])
x_gc = lm_gc.predict(train[['SalePrice']])
rmse_gc = mean_squared_error(train['GarageCars'], x_gc, squared=False)

plt.scatter(train['GarageCars'], train['SalePrice'], color='blue')
plt.plot(x_gc, train['SalePrice'], color='maroon', linewidth=1)
plt.ylabel('Sale Price ($)')
plt.xlabel('Size of garage (car capacity)')
plt.title('SalePrice vs GarageCars Regression')

```

```
print("RMSE:", rmse_gc)
```

RMSE: 0.573766659729406



The best predictors based on RMSE from lowest to highest are GarageCars (0.57), OverallQual (0.84), GrLivArea (370.64). GarageCars has the lowest RMSE which makes it the best predictor of the 3 tested above.

4. Engage in feature creation by splitting, merging, or otherwise generating a new predictor.

```
[23]: #Average Overall Condition and Quality
train['OverallCondQual'] = (train['OverallCond'] + train['OverallQual'])/2
```

This predictor may remove the need to include both Condition and Quality into a model if they are highly correlated.

```
[24]: # Combine total finished basement sq ft with above ground
train['total_FinSqft'] = train['TotalBsmtSF'] + train['GrLivArea']
```

This predictor may remove the need to include both above and below square footage if they are highly correlated.

```
[25]: # Years to sell since last remodel
train['YearsToSell'] = train['YrSold'] - train['YearRemodAdd']
```

This predictor may provide insight on how long it took to sell a property since its last remodel.

```
[26]: train.head()
```

```
[26]: MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
0         60      RL         65.0      8450   Pave   NaN      Reg
1         20      RL         80.0      9600   Pave   NaN      Reg
2         60      RL         68.0     11250   Pave   NaN      IR1
3         70      RL         60.0      9550   Pave   NaN      IR1
4         60      RL         84.0     14260   Pave   NaN      IR1

LandContour Utilities LotConfig ... MiscFeature MiscVal MoSold YrSold \
0         Lvl     AllPub    Inside ...         NaN         0         2    2008
1         Lvl     AllPub      FR2 ...         NaN         0         5    2007
2         Lvl     AllPub    Inside ...         NaN         0         9    2008
3         Lvl     AllPub    Corner ...         NaN         0         2    2006
4         Lvl     AllPub      FR2 ...         NaN         0        12    2008

SaleType SaleCondition SalePrice OverallCondQual total_FinSqft \
0         WD          Normal    208500             6.0         2566
1         WD          Normal    181500             7.0         2524
2         WD          Normal    223500             6.0         2706
3         WD      Abnorml    140000             6.0         2473
4         WD          Normal    250000             6.5         3343

YearsToSell
0           5
1          31
2           6
3          36
4           8
```

```
[5 rows x 83 columns]
```

5. Using the dependent variable, perform both min-max and standard scaling in Python.

```
[27]: #min-max scaling
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
train_minmax = train.copy()

numeric_cols = train_minmax._get_numeric_data().columns
train_minmax[numeric_cols] = scaler.fit_transform(train_minmax[numeric_cols])
```

```
train_minmax[numeric_cols].head()
```

```
[27]: MSSubClass  LotFrontage  LotArea  OverallQual  OverallCond  YearBuilt  \
0      0.235294    0.150685  0.033420    0.666667    0.500    0.949275
1      0.000000    0.202055  0.038795    0.555556    0.875    0.753623
2      0.235294    0.160959  0.046507    0.666667    0.500    0.934783
3      0.294118    0.133562  0.038561    0.666667    0.500    0.311594
4      0.235294    0.215753  0.060576    0.777778    0.500    0.927536
```

```
      YearRemodAdd  MasVnrArea  BsmtFinSF1  BsmtFinSF2  ...  3SsnPorch  \
0      0.883333    0.12250    0.125089    0.0    ...    0.0
1      0.433333    0.00000    0.173281    0.0    ...    0.0
2      0.866667    0.10125    0.086109    0.0    ...    0.0
3      0.333333    0.00000    0.038271    0.0    ...    0.0
4      0.833333    0.21875    0.116052    0.0    ...    0.0
```

```
      ScreenPorch  PoolArea  MiscVal    MoSold  YrSold  SalePrice  \
0      0.0      0.0      0.0  0.090909    0.50  0.241078
1      0.0      0.0      0.0  0.363636    0.25  0.203583
2      0.0      0.0      0.0  0.727273    0.50  0.261908
3      0.0      0.0      0.0  0.090909    0.00  0.145952
4      0.0      0.0      0.0  1.000000    0.50  0.298709
```

```
      OverallCondQual  total_FinSqft  YearsToSell
0      0.588235    0.195481    0.098361
1      0.705882    0.191802    0.524590
2      0.588235    0.207742    0.114754
3      0.588235    0.187336    0.606557
4      0.647059    0.263531    0.147541
```

[5 rows x 40 columns]

```
[28]: #standard scaling
from sklearn.preprocessing import StandardScaler
from scipy.stats import zscore
std_scaler = StandardScaler()
train_std = train.copy()

numeric_cols = train_std._get_numeric_data().columns
train_std[numeric_cols] = std_scaler.fit_transform(train_minmax[numeric_cols])
train_std[numeric_cols].head()
```

```
[28]: MSSubClass  LotFrontage  LotArea  OverallQual  OverallCond  YearBuilt  \
0      0.073375   -0.208034 -0.207142    0.651479   -0.517200    1.050994
1     -0.872563    0.409895 -0.091886   -0.071836    2.179628    0.156734
2      0.073375   -0.084449  0.073480    0.651479   -0.517200    0.984752
3      0.309859   -0.414011 -0.096897    0.651479   -0.517200   -1.863632
```

```
4      0.073375      0.574676  0.375148      1.374795      -0.517200      0.951632
```

```
      YearRemodAdd  MasVnrArea  BsmtFinSF1  BsmtFinSF2  ...  3SsnPorch  \  
0      0.878668      0.510015      0.575425      -0.288653  ...  -0.116339  
1     -0.429577     -0.572835      1.171992     -0.288653  ...  -0.116339  
2      0.830215      0.322174      0.092907     -0.288653  ...  -0.116339  
3     -0.720298     -0.572835     -0.499274     -0.288653  ...  -0.116339  
4      0.733308      1.360826      0.463568     -0.288653  ...  -0.116339
```

```
      ScreenPorch  PoolArea  MiscVal  MoSold  YrSold  SalePrice  \  
0     -0.270208 -0.068692 -0.087688 -1.599111  0.138777   0.347273  
1     -0.270208 -0.068692 -0.087688 -0.489110 -0.614439   0.007288  
2     -0.270208 -0.068692 -0.087688  0.990891  0.138777   0.536154  
3     -0.270208 -0.068692 -0.087688 -1.599111 -1.367655  -0.515281  
4     -0.270208 -0.068692 -0.087688  2.100892  0.138777   0.869843
```

```
      OverallCondQual  total_FinSqft  YearsToSell  
0      0.192175      -0.008372      -0.869941  
1      1.373547      -0.059386       0.390141  
2      0.192175      0.161672      -0.821476  
3      0.192175      -0.121330       0.632464  
4      0.782861      0.935372      -0.724547
```

```
[5 rows x 40 columns]
```

When using machine learning, models are typically used to tell the distance between data. By scaling using either min-max or standard, we can normalize and/or standardize the data before going into certain analyses or algorithms that require it.

```
[ ]:
```